

ETHIC VAULT: EXECUTIVE BRIEF

A Proposed Institutional Framework for Rewarding Social Coherence

Document Type: Executive Brief (Social Sciences Edition)

Target Audience: Economists, Sociologists, Behavioral Scientists

Based on: Ethic Vault Thought Paper on the Dual Closed-Loop Economy

Purpose: Invitation to Collaborative Experimental Design

Author: George Dunphy, Ethic Vault

Date: May 2026

Status: Preprint for Academic Collaboration

1. THE PROBLEM: A MARKET FAILURE HIDDEN IN PLAIN SIGHT

Digital platforms have solved many coordination problems, but they have simultaneously created a new class of market failure: the systematic underproduction of social public goods.

Trust, mentorship, dispute resolution, fraud prevention, and governance participation are not luxuries; they are the fundamental infrastructure of cooperation. Yet, in the current architecture of the internet, these activities remain invisible labor—socially necessary, economically unrewarded, and therefore systematically undersupplied.

The mechanisms that do exist are misaligned. Platforms today do not reward constructive contribution; they reward emotional volatility. Outrage generates attention, and attention generates revenue. This is not a moral failure of individuals; it is a design failure of incentives.

Three well-documented phenomena illustrate the depth of this problem:

Goodhart's Law in Action: When platforms measure engagement, users and algorithms optimize strictly for engagement. The metric eats the value it was meant to proxy. Helpfulness, accuracy, and trust are not measured, so they are not produced.

Behavioral Rent-Seeking: Platforms extract value by amplifying cognitive biases—outrage, fear, tribal loyalty—because these are reliably monetizable. The resulting social costs (polarization, harassment, misinformation) are treated as unpriced externalities.

The Multi-Task Principal-Agent Problem: Organizations and platforms want helpers, mentors, and dispute resolvers, but they reward only what is easy to measure: volume, speed, and visibility. The invisible work of de-escalation, quality review, and behind-the-scenes governance is systematically devalued.

The consequence is not merely inefficiency; it is institutional fragility. When trust fails, coordination costs rise. When coordination costs rise, cooperation collapses.

2. THE PROPOSED MECHANISM: A BEHAVIORAL-INSTITUTIONAL SOLUTION

We propose a mechanism designed specifically to address this market failure: a dual closed-loop economy. The name is descriptive, not technical.

Loop One represents conventional economic production: work, transactions, services, and capital formation. This loop already has established measurement and reward systems. They are imperfect, but they exist.

Loop Two is what we propose to formalize: a mechanism for the verification, attribution, and rewarding of pro-social behavior. This includes verified help, mentorship, dispute resolution, fraud reporting, governance participation, and any activity that lowers coordination costs and raises system trust.

The two loops are not separate; they are coupled by design. Loop One funds Loop Two through budgets and access policies, while Loop Two feeds back into Loop One through lower fraud rates, faster dispute resolution, better partner selection, and higher-quality decision inputs. Social coherence stops being invisible labor and becomes a priced input to institutional performance.

What makes the system a "closed-loop" is an economic feature rather than a technical one: evidence, rewards, sanctions, and governance outcomes flow back into future access, pricing, and allocation decisions rather than disappearing as unpriced externalities. This is the solution to the tragedy of the commons applied to social behavior: operational feedback creates structural accountability.

Furthermore, the mechanism directly addresses the multi-task principal-agent problem by rewarding multiple dimensions of pro-social contribution through a bundle of signals rather than a single scalar score. Helpfulness, uniqueness, evidence quality, and outcome significance are all weighted organically so that no single metric can dominate or distort the system.

3. TESTABLE RESEARCH QUESTIONS: A COLLABORATION MENU

We do not claim to have solved these problems; we claim to have built a mechanism that can empirically test hypotheses about them. We invite collaboration on the following research questions:

Q1. Non-transferable credits and cooperation.

How does the introduction of a non-transferable "coherence credit"—earned through verified pro-social behavior but not convertible to currency—affect rates of cooperation, altruism, and contribution quality over time? Does the absence of liquidity change the social meaning of the reward?

Q2. Separation of capital and social capital.

Does separating reputation from financial liquidity prevent the unfair conversion of economic capital into social capital (and vice versa)? Under what conditions does this separation break down? Can token-weighted governance be made capture-resistant without sacrificing legitimacy?

Q3. Procedural justice in automated governance.

What defines "procedural justice" in a digital system where sanctions, credential issuance, and appeals are partially automated? Under what conditions do human participants accept machine-led governance decisions? When do they demand human review? What is the minimal viable appeals process for perceived legitimacy?

Q4. Resilience to manipulation.

How do multivariate reward functions, delayed finalization, random audits, and pairwise earning caps perform against adversarial gaming relative to single-score systems? Can specification gaming be systematically suppressed through mechanism design, or will optimizers always find a loophole?

Q5. Social welfare effects.

Does making social coherence economically visible and rewarded produce measurable improvements in allocative efficiency, fraud reduction, dispute closure speed, or participant retention? What are the unintended consequences?

4. THE INVITATION: CO-DESIGN A CONTROLLED SOCIAL EXPERIMENT

We are not launching a finished product; we are designing a controlled social experiment.

The technology stack exists for one primary reason: to ensure the auditability and integrity of the experiment. It is the infrastructure that allows us to run the experiment without losing data to tampering, without losing identity to Sybil attacks, and without losing trust to opaque governance.

We invite economists, sociologists, and behavioral scientists to co-design this experiment with us, specifically collaborating on:

- Dependent variable selection: What are the right measures of trust, coherence, and allocative quality? How do we operationalize them for empirical testing?
- Experimental design: Should the first pilots use randomization, control groups, or staggered rollouts? How do we isolate the effect of the mechanism from confounding variables?
- Measurement protocols: How do we measure Goodhart resistance empirically? What constitutes definitive evidence of specification gaming?
- Qualitative methods: How do we capture participant perceptions of fairness, legitimacy, and procedural justice alongside quantitative metrics?
- Comparative analysis: How does this mechanism compare to existing trust systems (e.g., Stack Overflow, Community Notes, Gitcoin grants) on measurable outcomes?

We have identified three initial pilot environments where social coherence already has visible value but poor measurement:

1. Professional support or learning communities (focusing on mentorship and verified help)
2. Enterprise support or fraud prevention workflows (focusing on dispute resolution and fraud reporting)

3. Grants or open-source communities (focusing on contribution attribution and collective allocation)

Each pilot is an ideal candidate for rigorous, pre-registered experimental evaluation with academic oversight.

5. ETHICAL SAFEGUARDS: WHAT THIS IS NOT

Because any system that measures and rewards social behavior risks being misunderstood, we state clearly: This is NOT a social credit system.

The differences are substantive, not semantic:

Feature	Social Credit Systems	Ethic Vault Framework
Participation	Coercive, state-mandated	Opt-in, context-specific
Direction	Punishes deviation from norms	Rewards positive contribution
Scope	Universal, totalizing	Bounded to specific platforms/communities
Transparency	Black box, non-appealable	Policy-as-code, full auditability, appeals
Data collection	Maximum surveillance	Minimal disclosure, off-chain PII

We take the risk of mission drift seriously. The governance charter for any pilot will include explicit prohibitions on:

- Using coherence credits to deny essential services
- Punishing off-platform behavior
- Operating without transparent appeal and correction rights
- Allowing any single actor to control the reward function
- Expanding scope beyond the agreed pilot boundaries without explicit re-consent

If the system ever feels like coercive social scoring, it has already failed its own governance test. That is not a possibility we dismiss; it is a risk we design against from first principles.

6. A REQUEST FOR DIALOGUE

We do not need validation; we need collaboration.

The technical infrastructure is ready for small-scale, adversarial, and transparent pilots. What is missing is the rigorous social science framing that will allow us to interpret results, identify confounders, and iterate toward a mechanism that actually works—not just technically, but behaviorally and institutionally.

We invite you to:

- Review the full technical paper (available on request)
- Suggest modifications to the experimental design

- Propose alternative dependent variables or measurement strategies
- Identify risks we have not anticipated
- Join a pilot as a design partner or independent evaluator
- Help us distinguish genuine social coherence from performative behavior
- Design the qualitative research protocols alongside the quantitative ones

The internet already has an entrenched reward function, and it is actively producing outrage, polarization, and manipulability. We are proposing a testable alternative. The question is not whether it is perfect; the question is whether, with your help, it can be better.

We are not asking for endorsement. We are asking for rigorous, critical, and constructive engagement. If the mechanism is flawed, we want to discover that early, transparently, and with academic rigor. If it shows promise, we want to understand why, under what conditions, and with what limitations.

7. NEXT STEPS FOR COLLABORATION

For those interested in engaging with this work, we propose the following pathway:

Step 1: Initial conversation (30-45 minutes)

- Walk through the mechanism design assumptions
- Identify immediate concerns or red flags
- Discuss fit with your research interests

Step 2: Experimental design workshop (half-day, virtual or in-person)

- Co-design the dependent variables and measurement protocols
- Review the proposed pilot environments
- Identify potential confounders and control strategies

Step 3: Pilot observation or advisory role

- Observe the first pilot cycle (approximately 8-12 weeks)
- Provide feedback on observed behavior and outcomes
- Suggest mid-course corrections

Step 4: Co-authorship of findings

- Jointly analyze the results
- Co-author working papers or journal submissions
- Present findings at relevant conferences

All collaboration can be anonymous or attributed, as preferred. We are equally open to formal research partnerships and informal advisory relationships.

8. CONCLUSION

The dual closed-loop economy is not a prediction; it is a proposal. It is not a product; it is an experiment.

We have built the infrastructure. We have grounded it in mechanism design, Ostrom's institutional principles, and the empirical reality of Goodhart's Law and specification gaming. What we need now is the social science expertise to test whether it actually works—not in theory, but in practice, with real humans, real incentives, and real adversarial pressure.

The question we put before you is simple: Will you help us test this?

We are not promising solutions. We are promising a rigorous, transparent, and auditable experimental platform. What we learn from it—whether success or failure—will be shared openly. The goal is not to prove that we were right; the goal is to build evidence about what actually produces social coherence in digital environments.

That is a question worth answering together.

Contact: George Dunphy, Ethic Vault

Email: [insert email]

Based on: Ethic Vault Thought Paper on the Dual Closed-Loop Economy

Date: May 2026